

emcee: An Affine-Invariant Sampler

Dustin Lang
Perimeter Institute for Theoretical Physics

PSI Numerical Methods, 2026-01-23

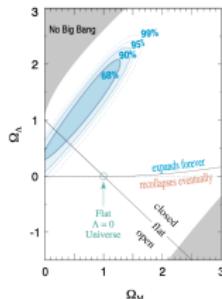
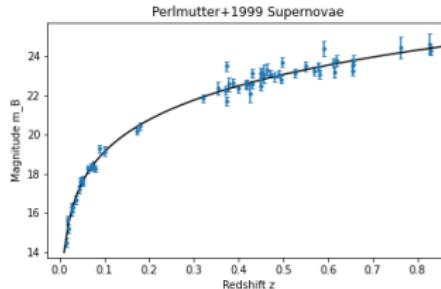
Borrowing heavily from Dan Foreman-Mackey's slides
<https://speakerdeck.com/dfm/data-analysis-with-mcmc1>

These slides are available at

<https://github.com/dstndstn/MCMC-talk/emcee-slides>

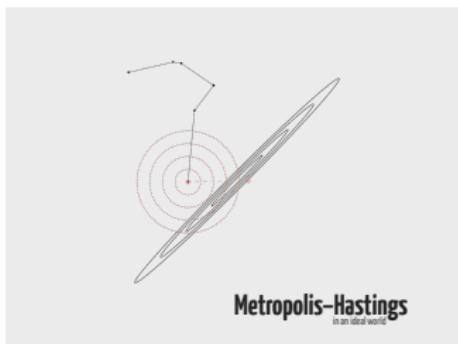
Recap from last lecture (1)

- ▶ Markov Chain Monte Carlo (MCMC) *draws samples from a probability distribution* when you can *numerically evaluate* the probability function (up to a constant)
- ▶ Used extensively in data analysis: *inferring* parameters of models, given observed data
- ▶ *Usually* in a Bayesian context; the probability function we run MCMC on is the *posterior* probability:
 $\text{posterior}(\text{params}|\text{data}) \propto \text{prior}(\text{params}) \times \text{likelihood}(\text{data}|\text{params})$

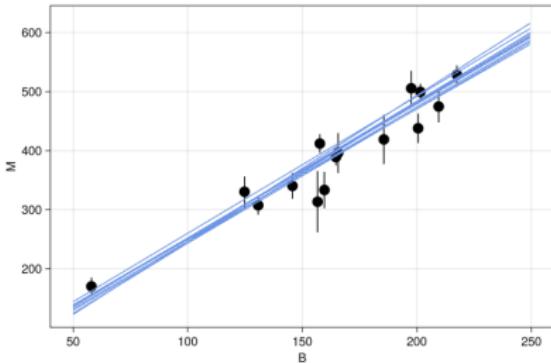
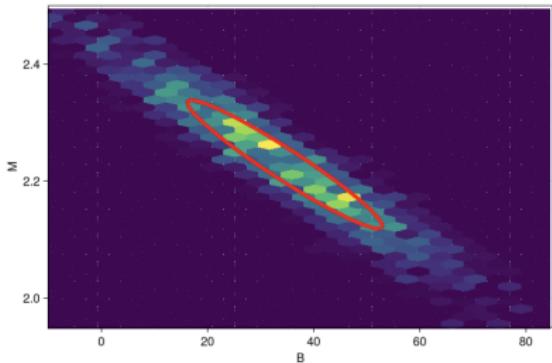
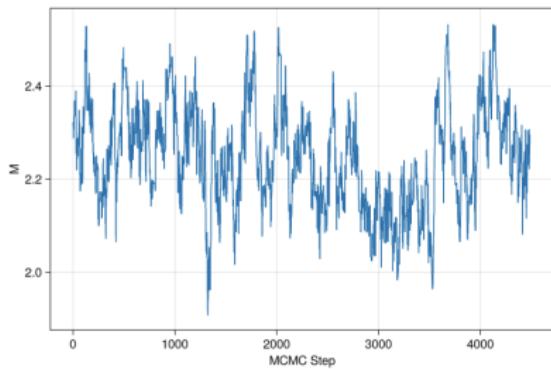
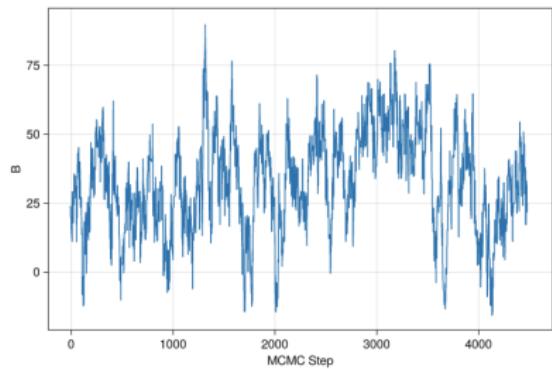


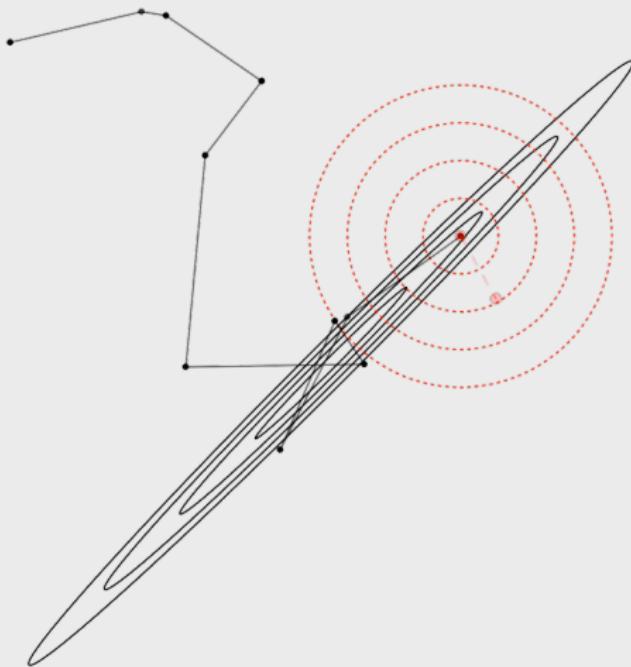
Recap from last lecture (2)

- ▶ The “classic” Markov Chain Monte Carlo algorithm is *Metropolis–Hastings*, which moves a *walker* or *particle* around the *state space* (*model parameter space*)
- ▶ A randomly-drawn *proposed* jump gets *evaluated* (by calling the probability function), and then *accepted*, or not
- ▶ A big difficulty is to *customize* the *proposal distribution* to get the algorithm to work efficiently



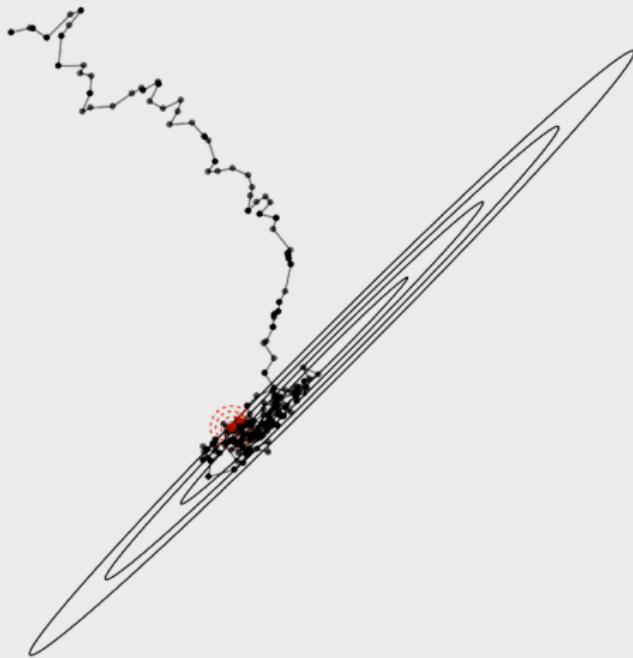
MCMC for model parameter inference





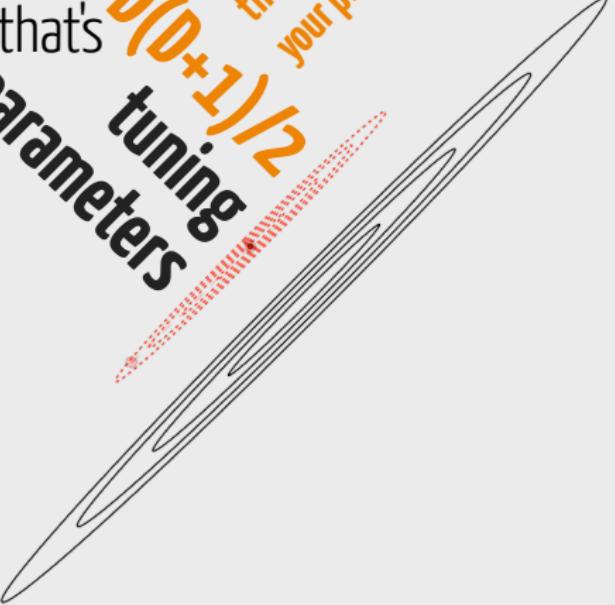
Metropolis–Hastings

in the real world



Metropolis–Hastings

in the real world



that's $D(D+1)/2$ tuning parameters
the dimension of your problem

Metropolis–Hastings

in the real world



Jonathan Goodman



Jonathan Weare

"Ensemble samplers with affine invariance"

(dfm.io/mcmc-gw10)



introducing
emcee the MCMC Hammer

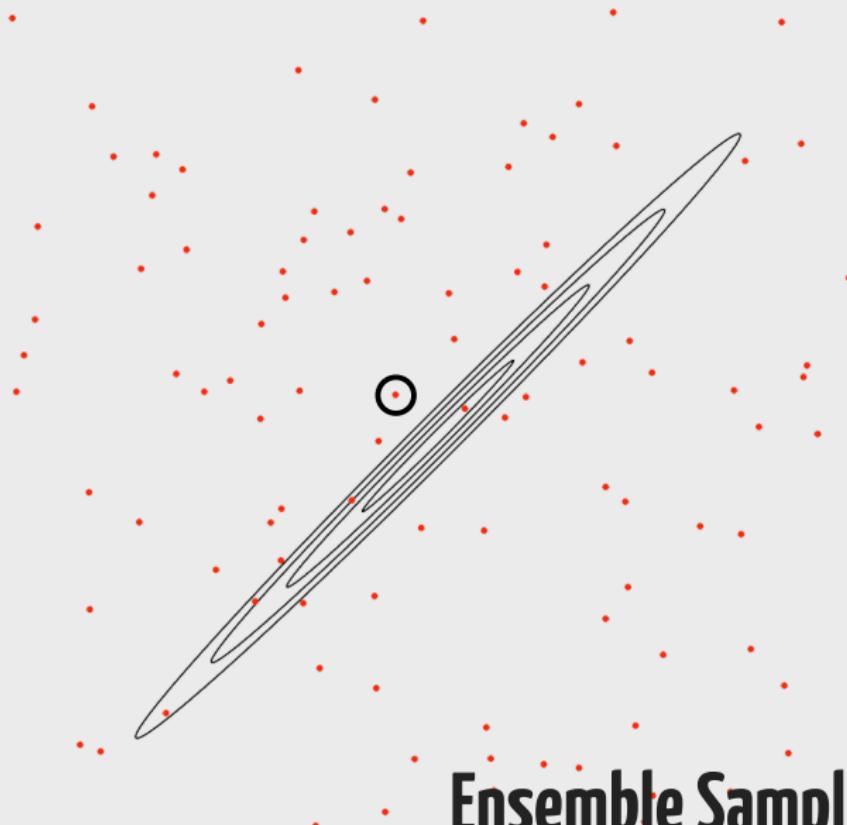
arxiv.org/abs/1202.3665

dan.iel.fm/emcee



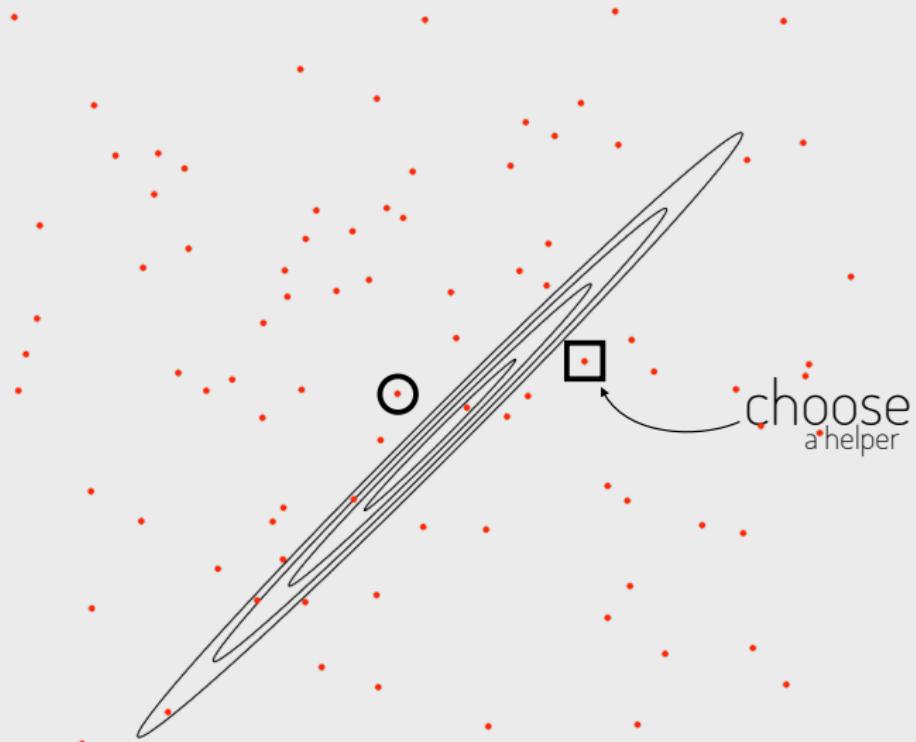
Ensemble Samplers

in the real world



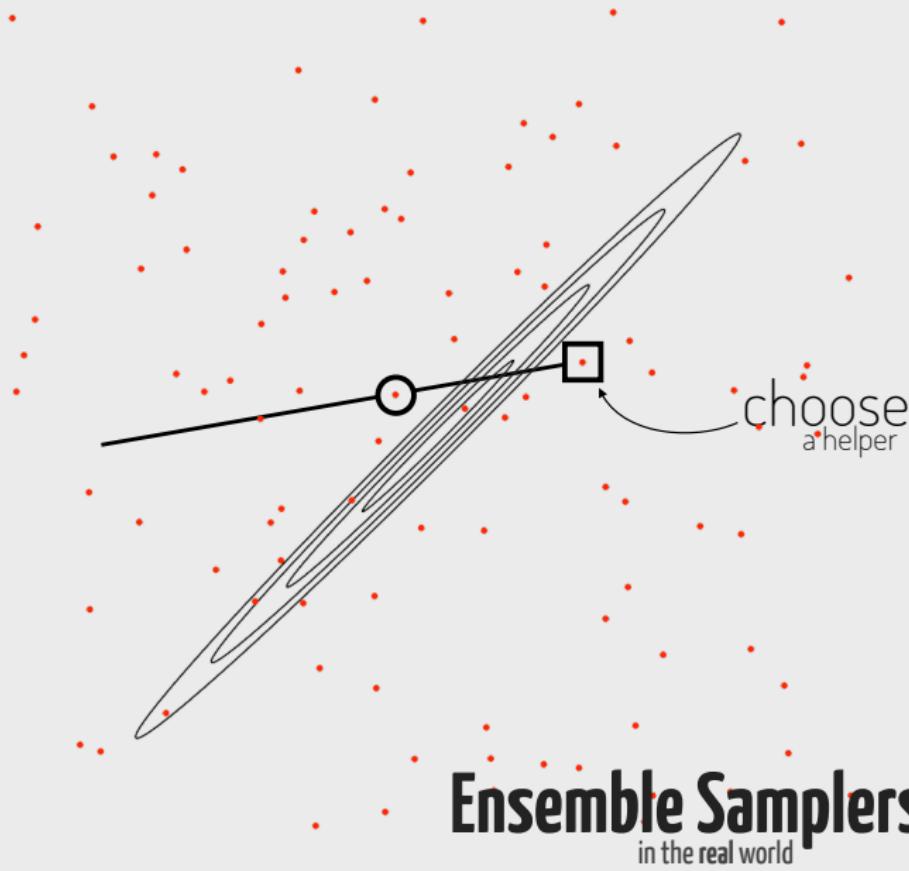
Ensemble Samplers

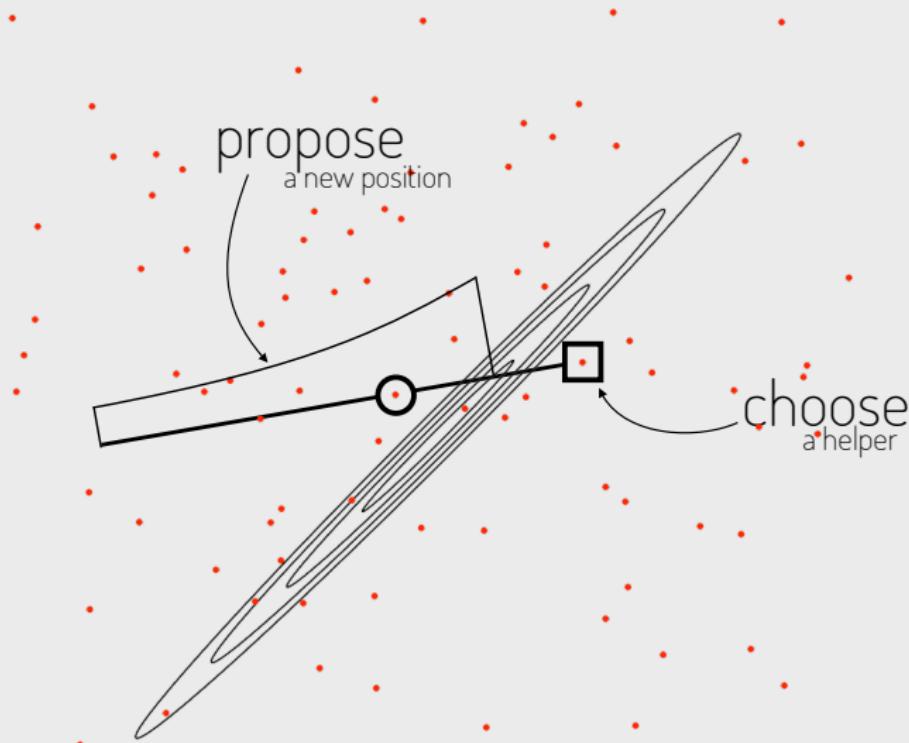
in the real world



Ensemble Samplers

in the real world

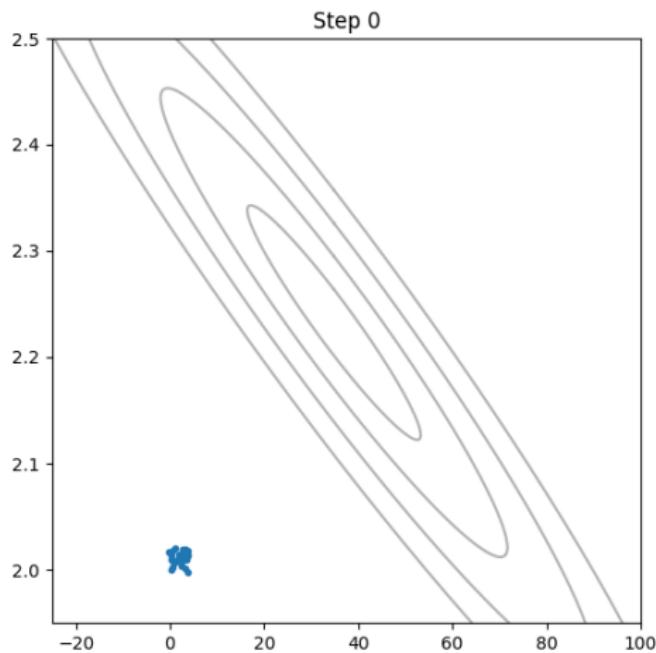




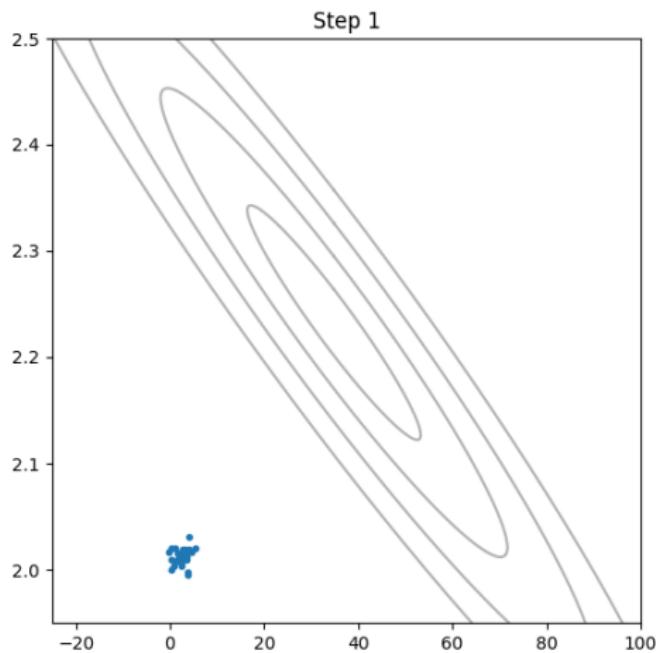
Ensemble Samplers

in the real world

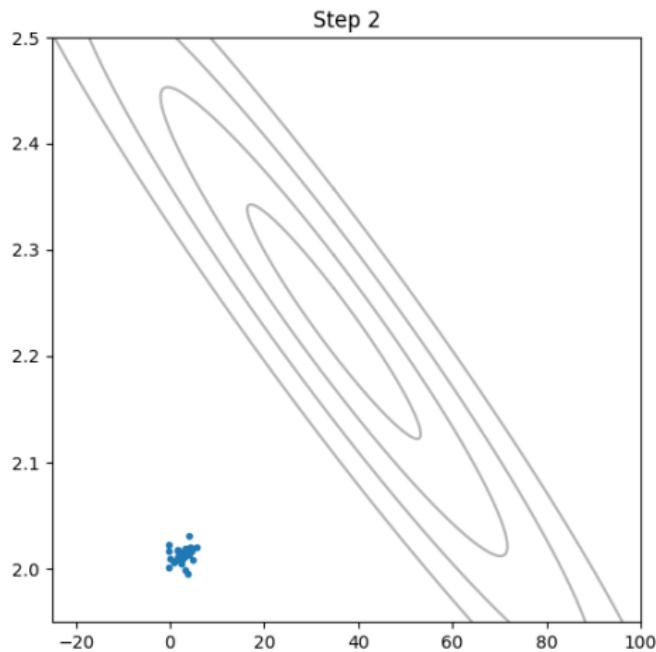
Emcee demo



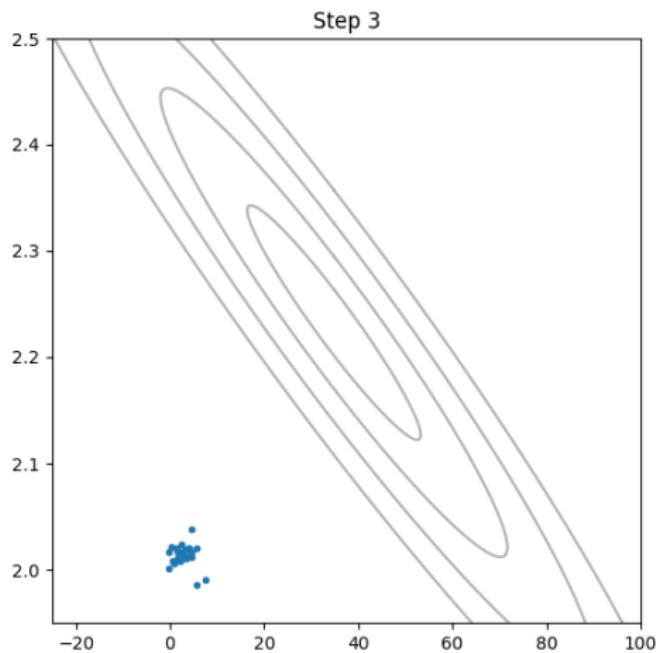
Emcee demo



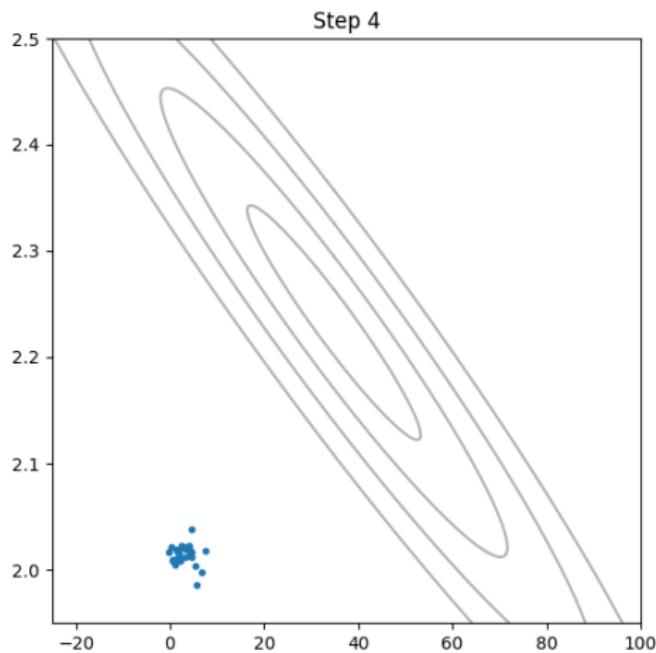
Emcee demo



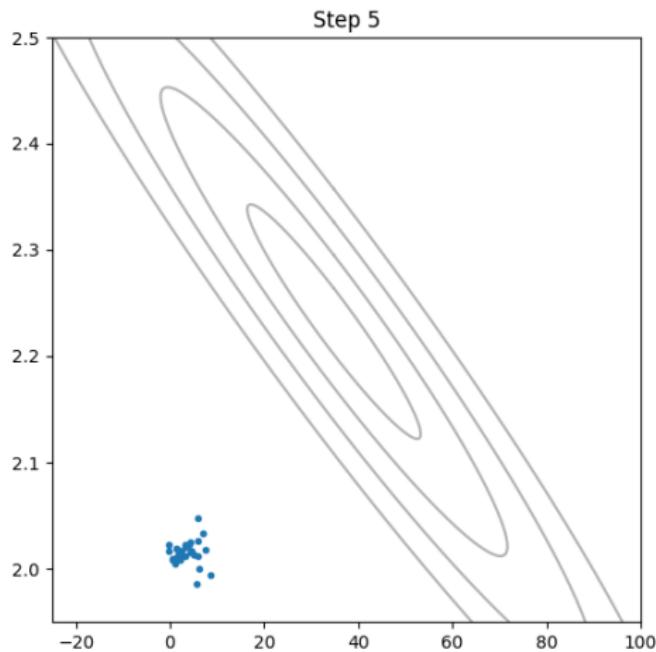
Emcee demo



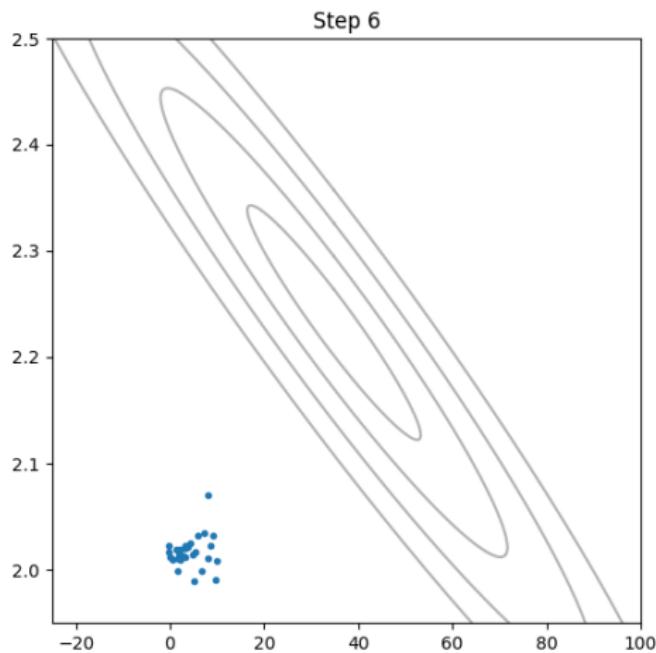
Emcee demo



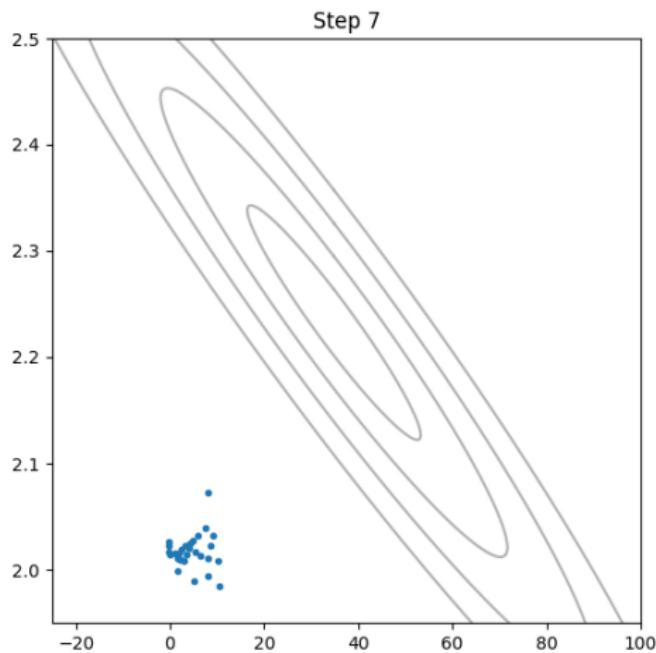
Emcee demo



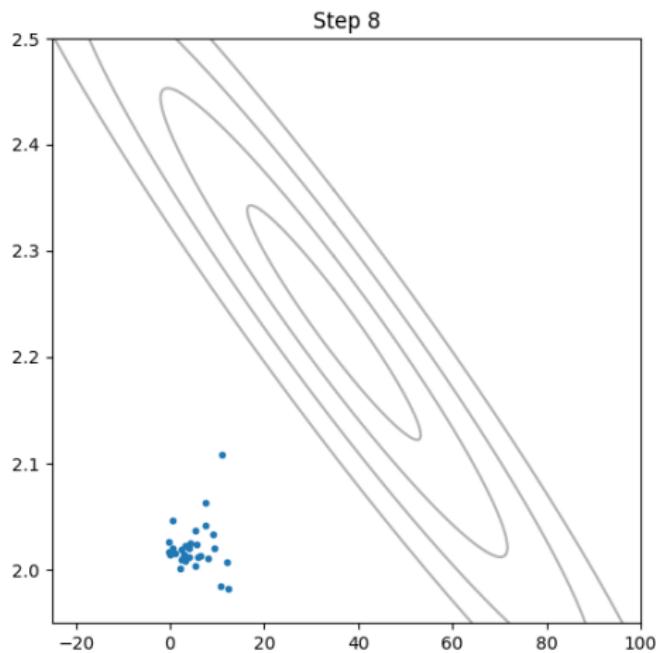
Emcee demo



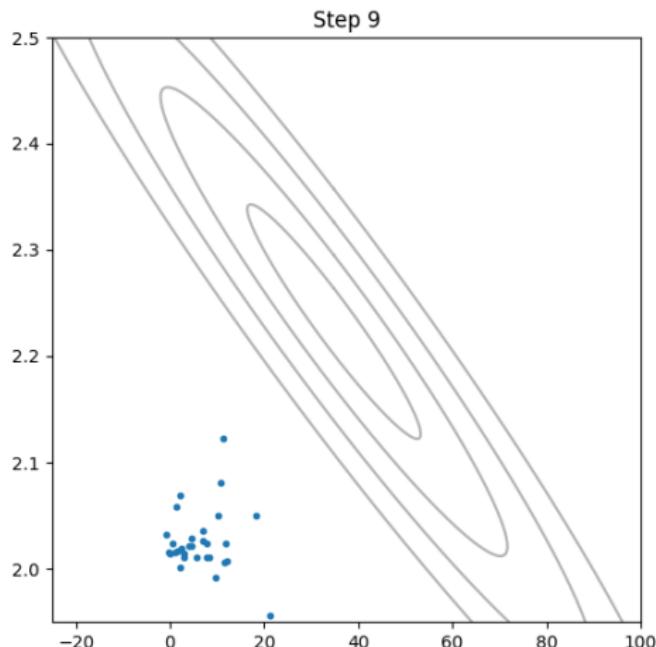
Emcee demo



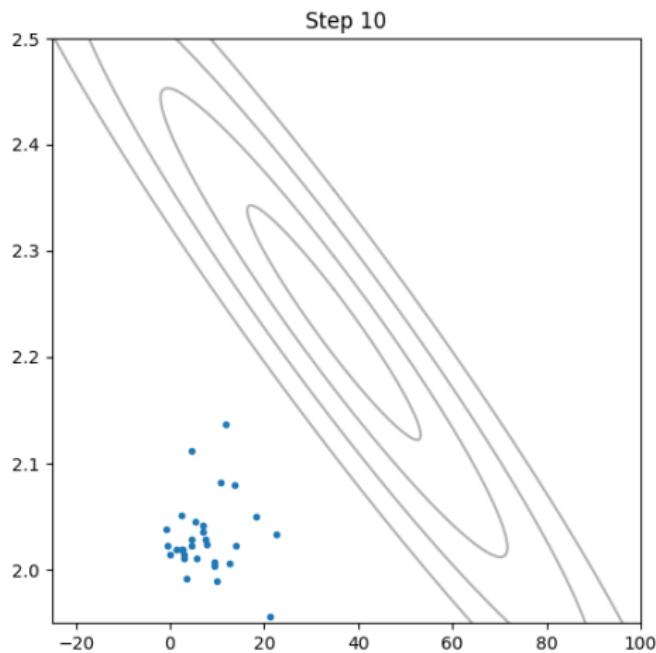
Emcee demo



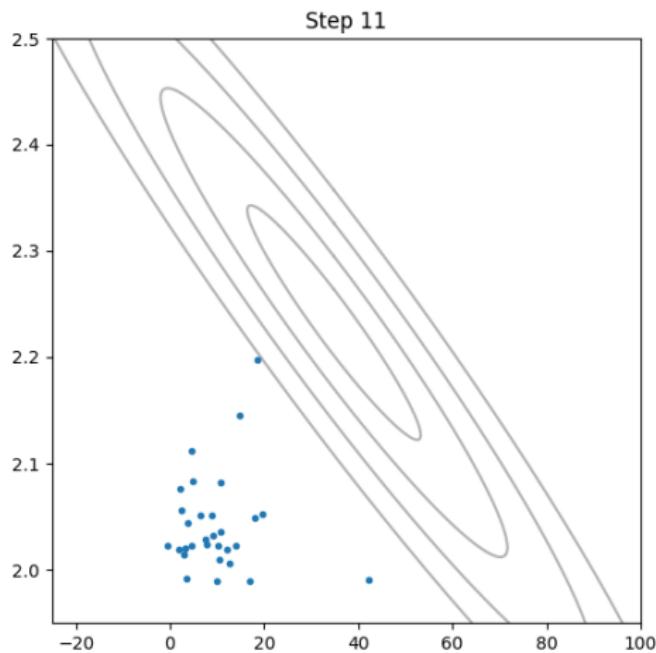
Emcee demo



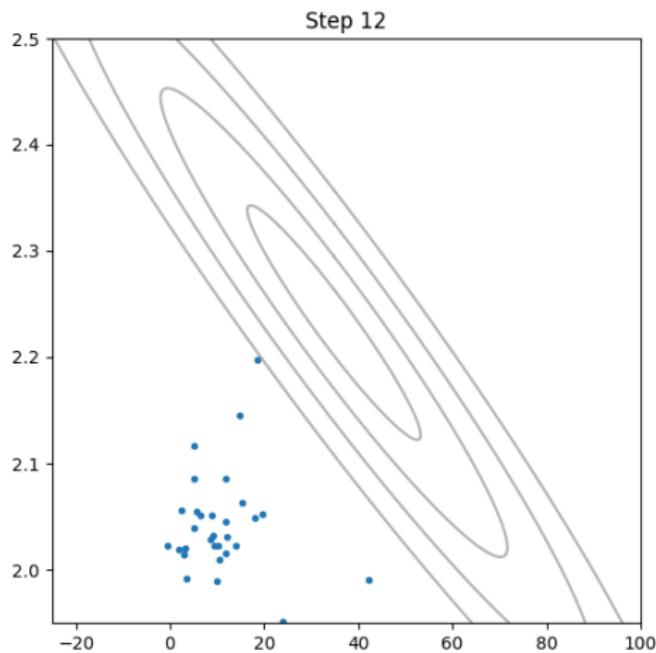
Emcee demo



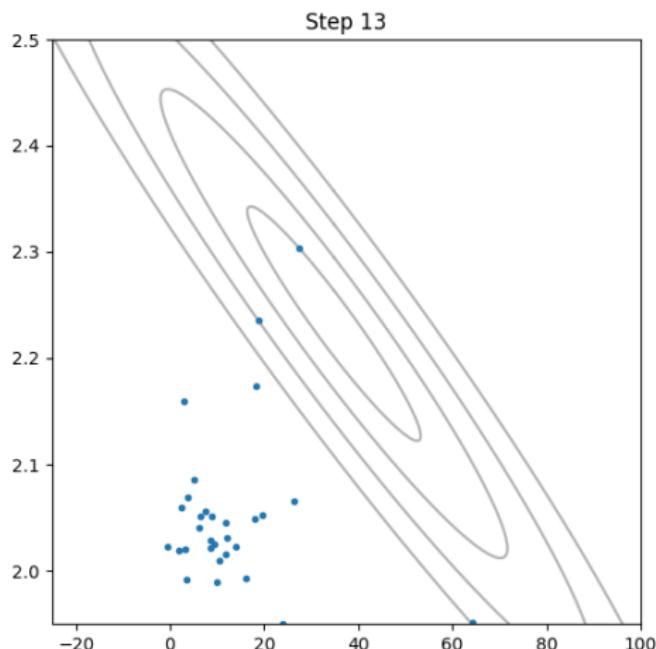
Emcee demo



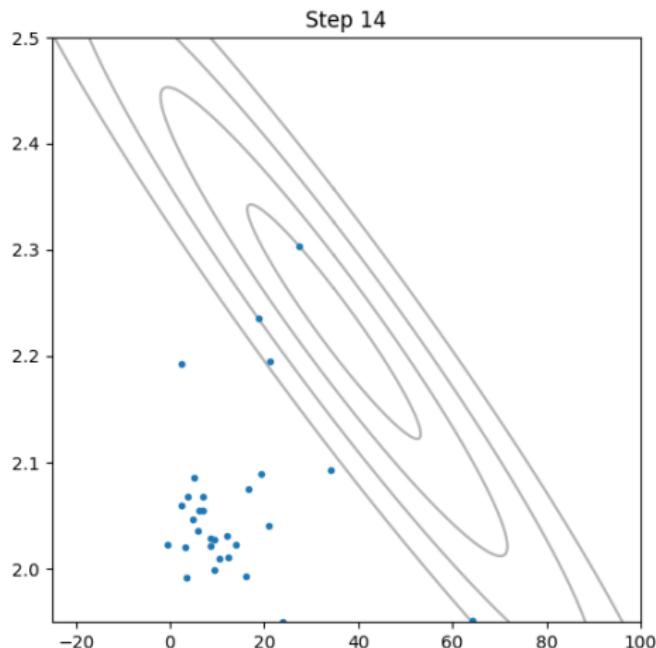
Emcee demo



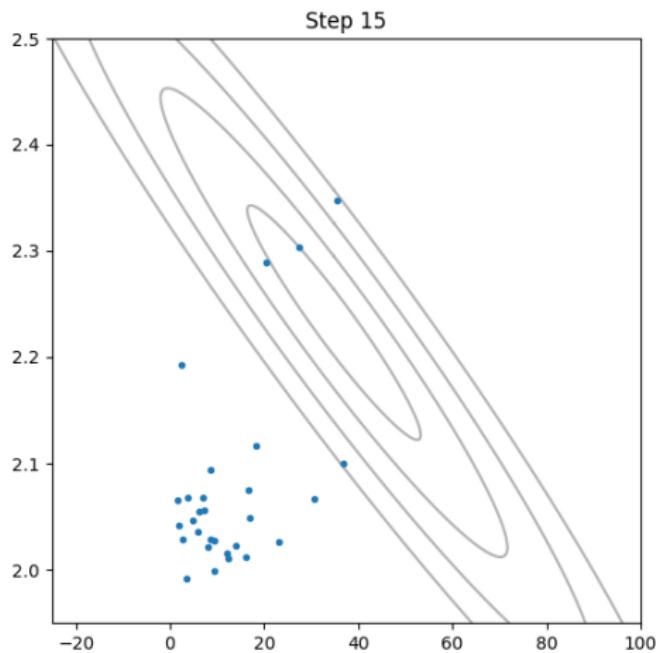
Emcee demo



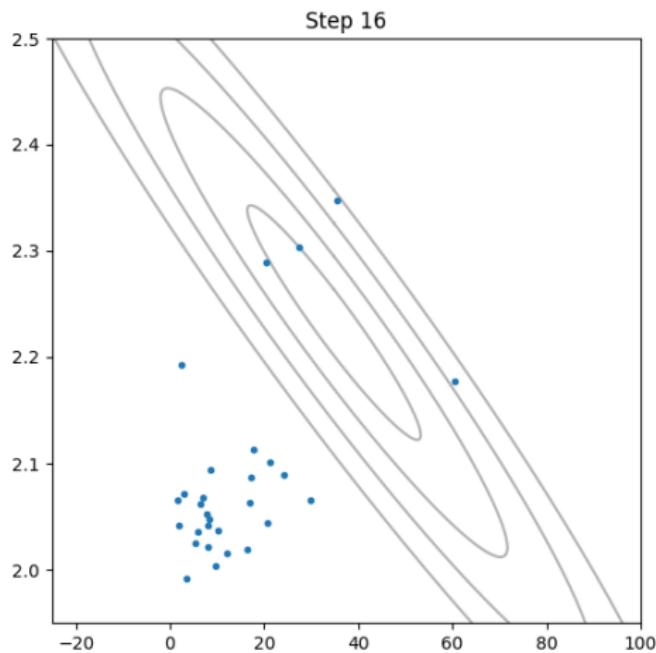
Emcee demo



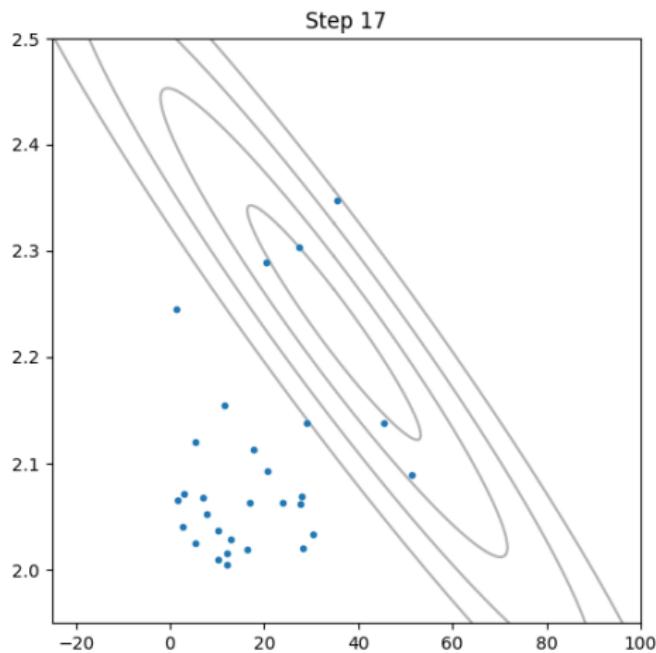
Emcee demo



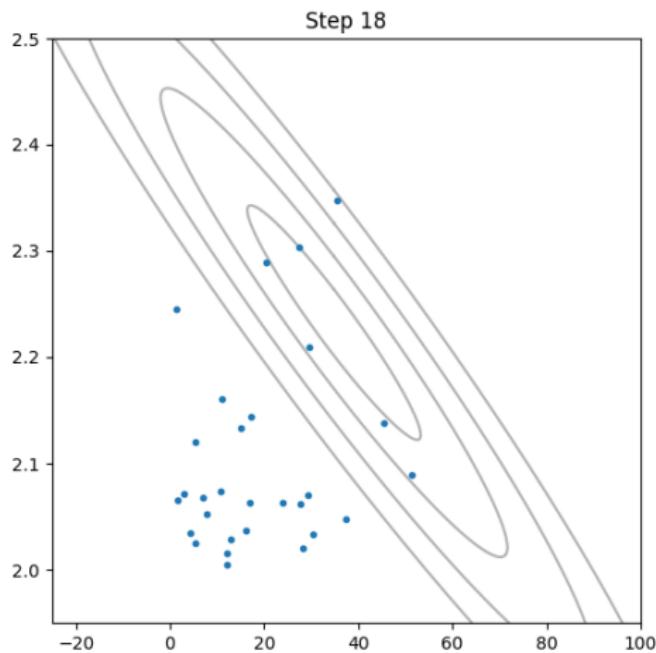
Emcee demo



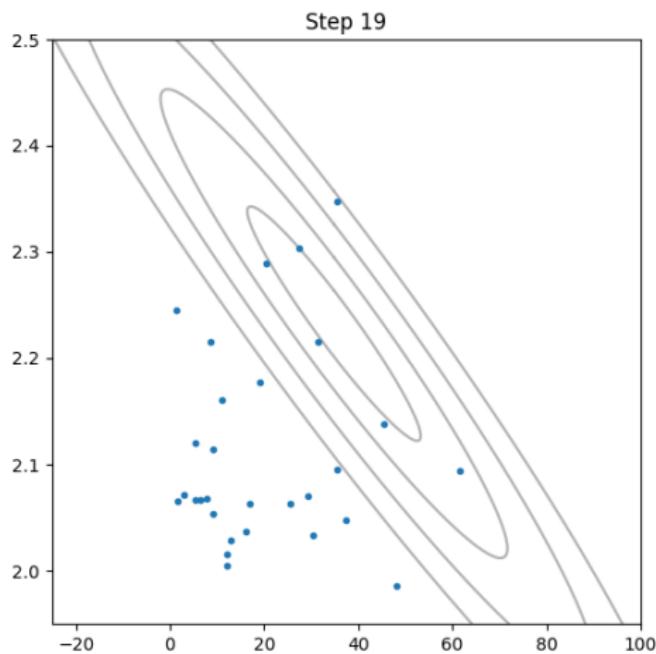
Emcee demo



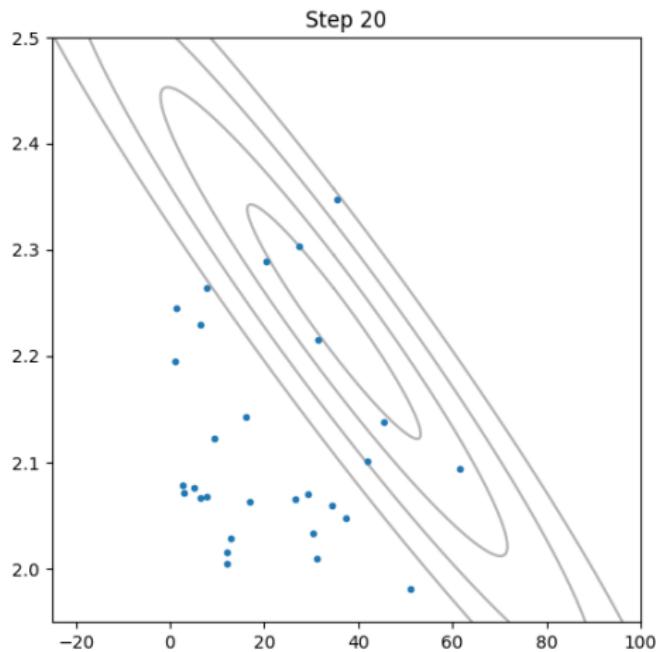
Emcee demo



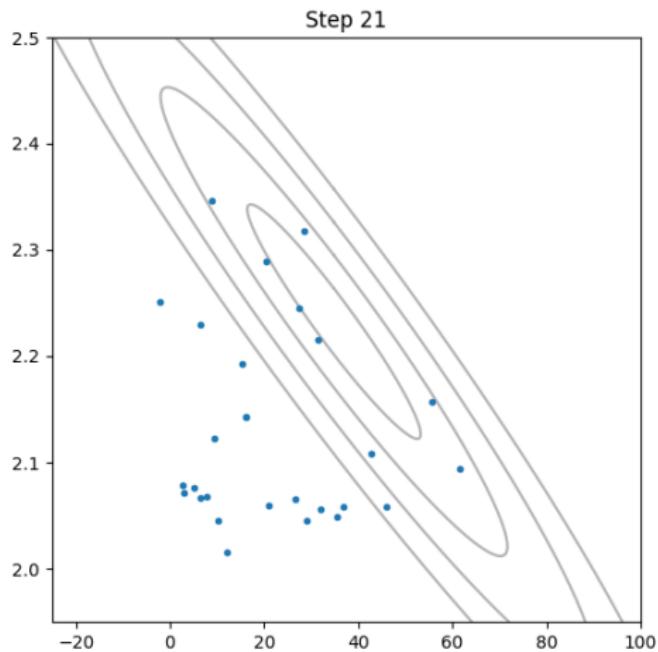
Emcee demo



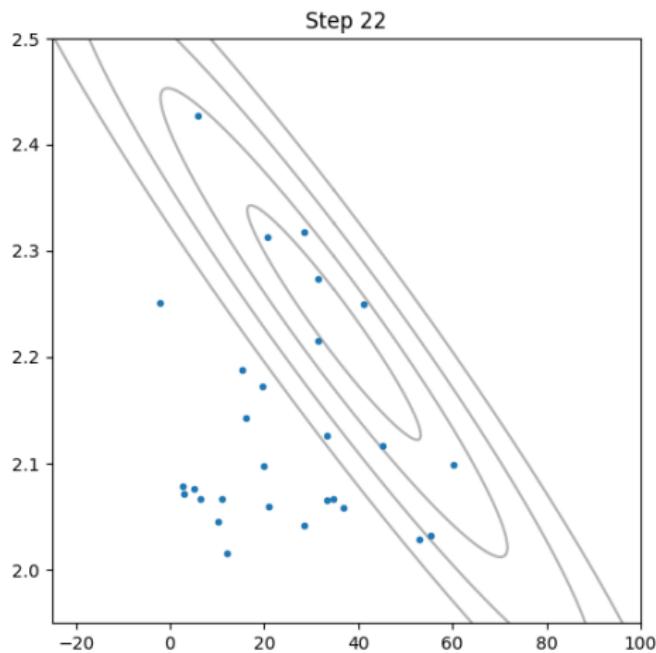
Emcee demo



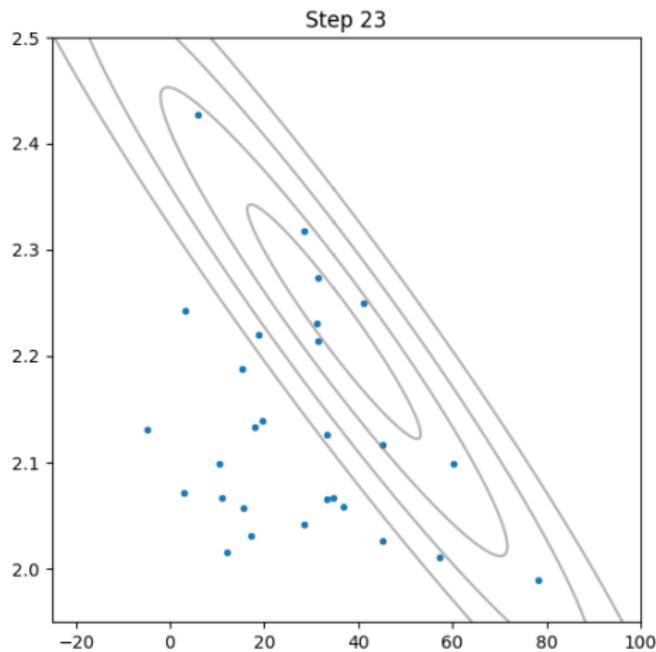
Emcee demo



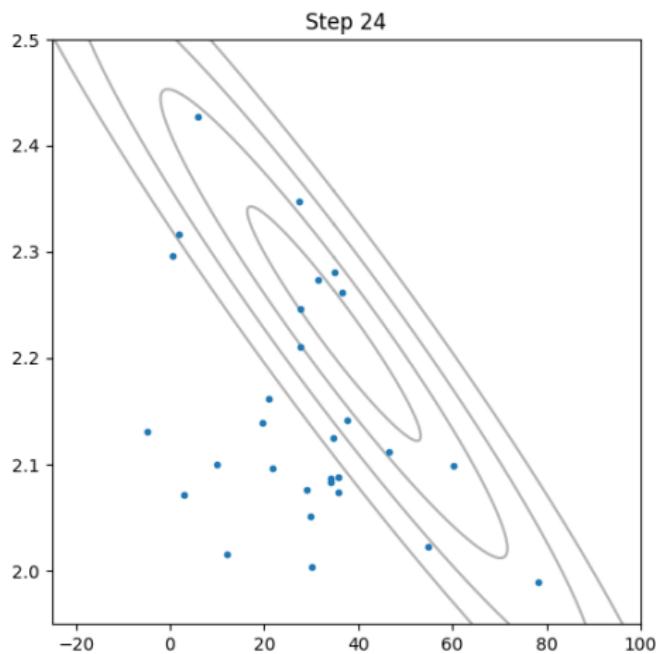
Emcee demo



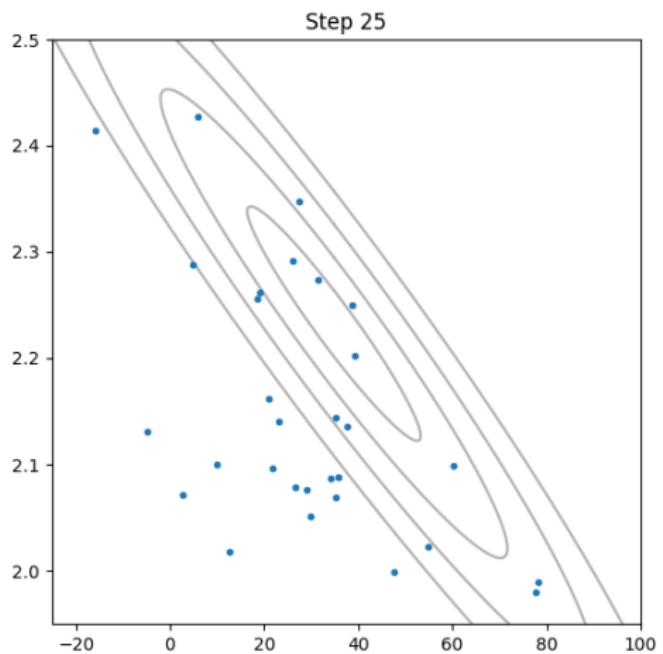
Emcee demo



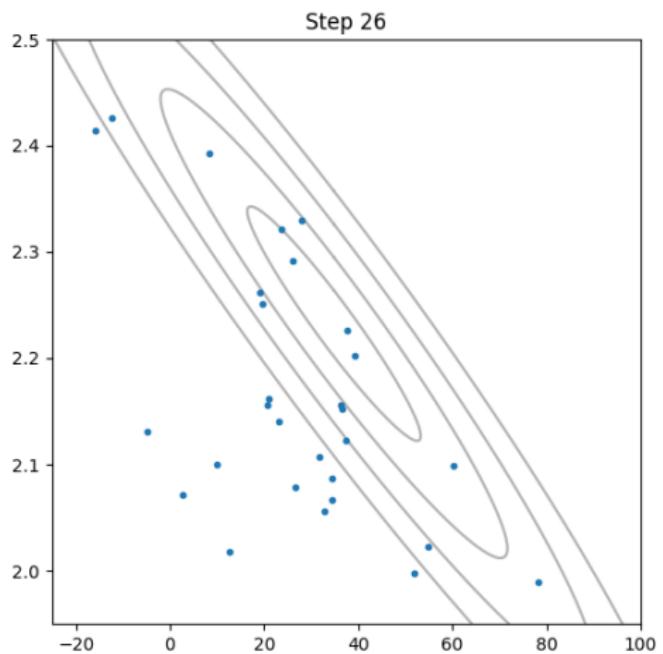
Emcee demo



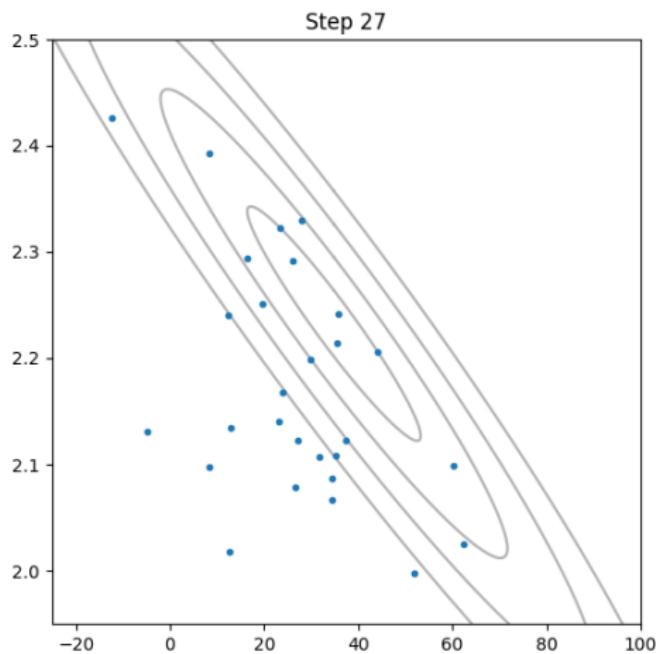
Emcee demo



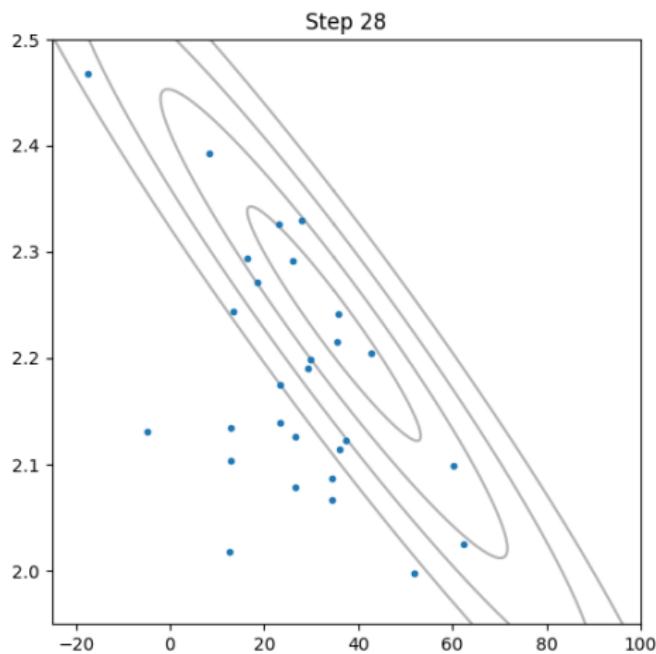
Emcee demo



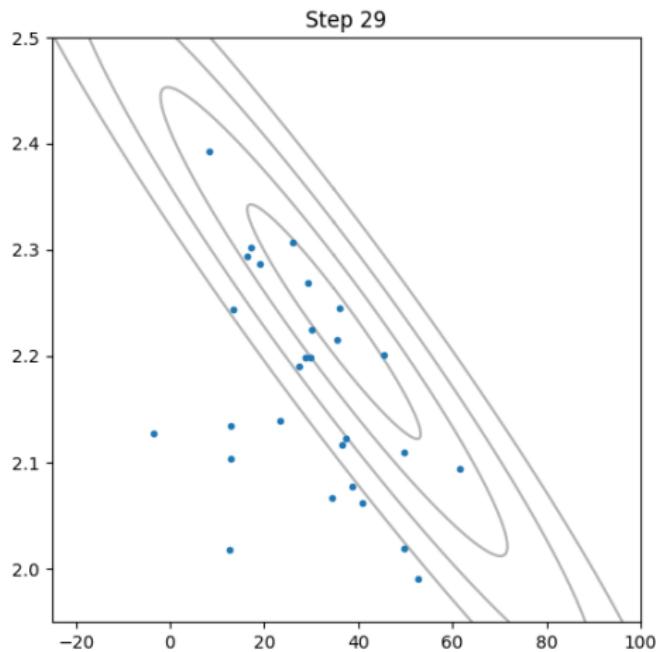
Emcee demo



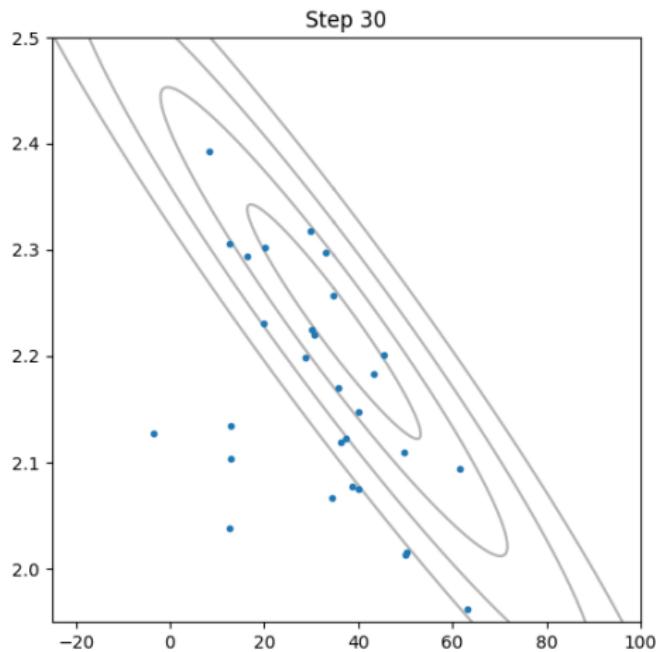
Emcee demo



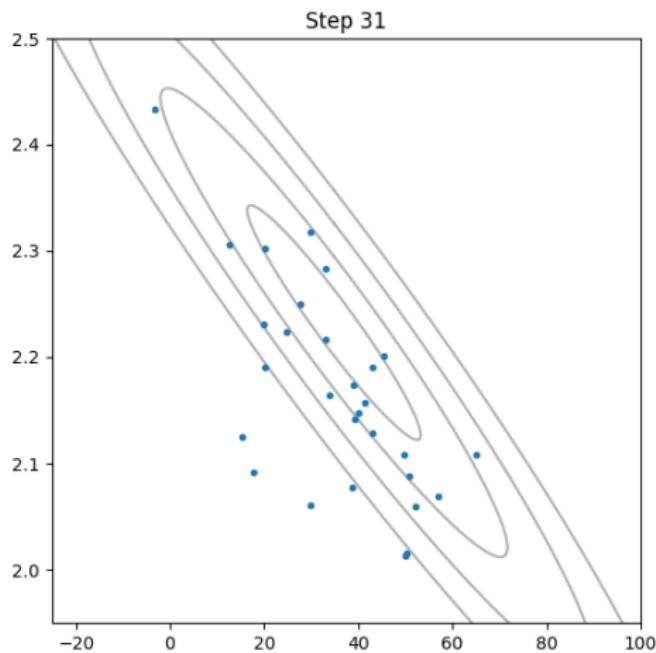
Emcee demo



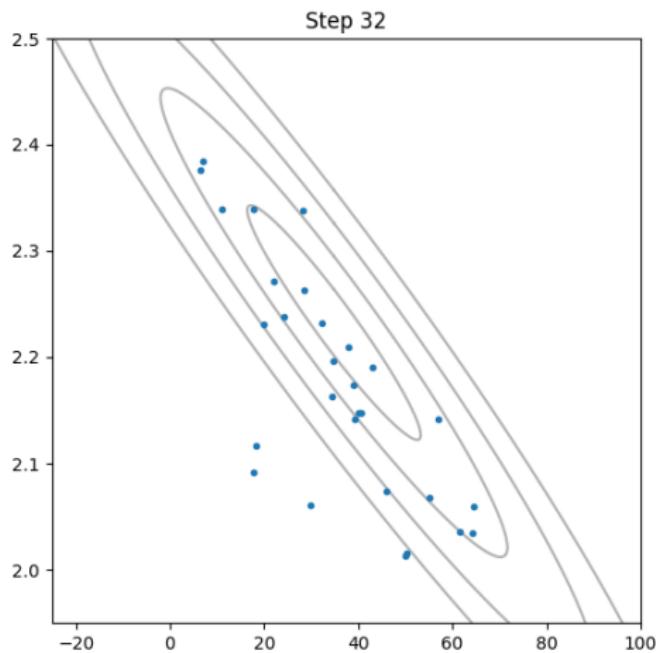
Emcee demo



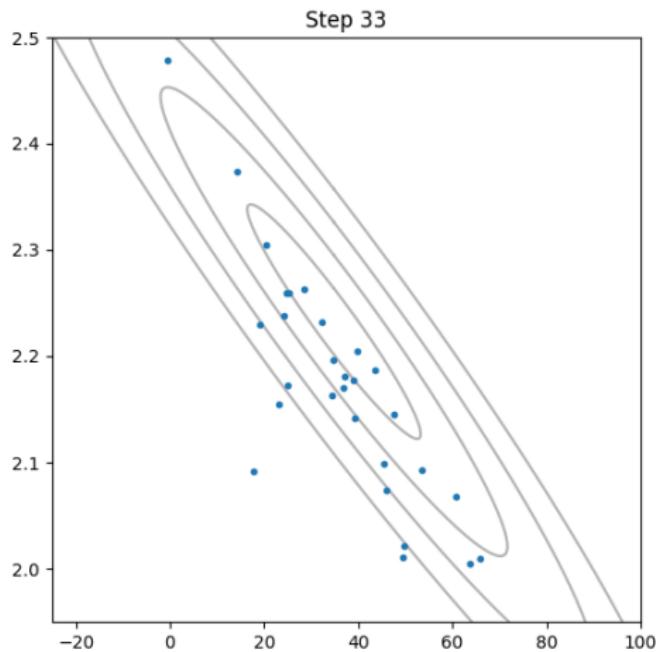
Emcee demo



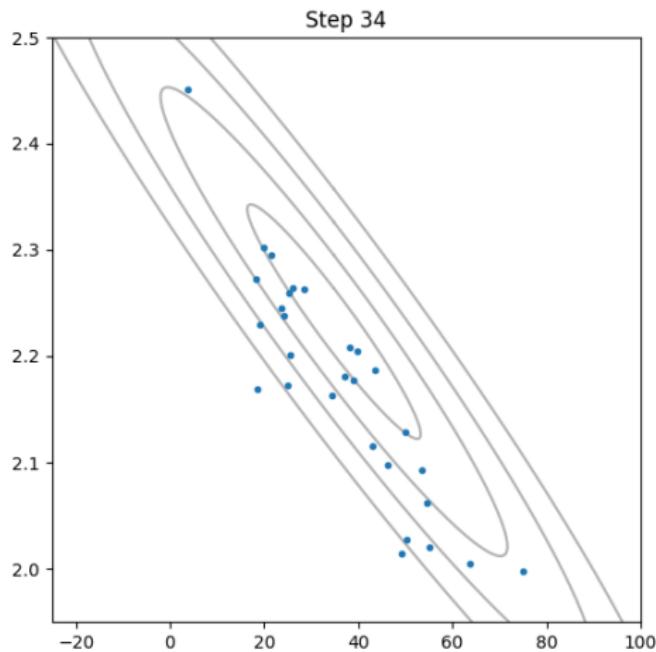
Emcee demo



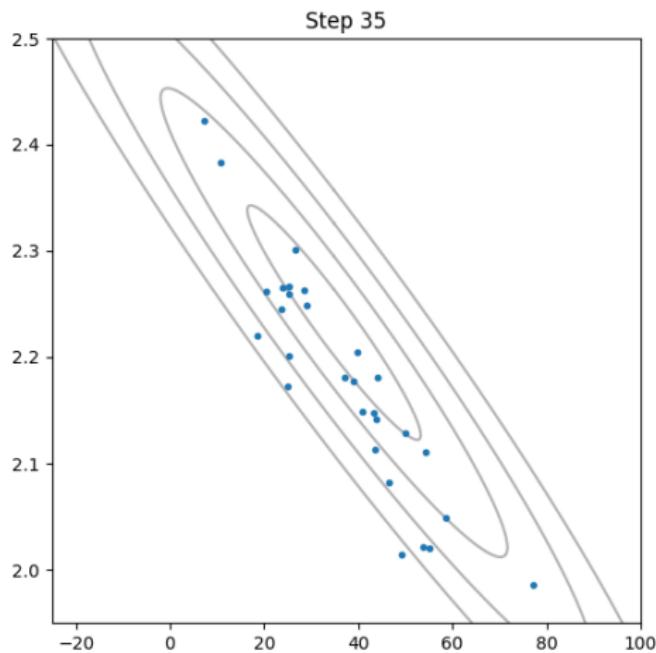
Emcee demo



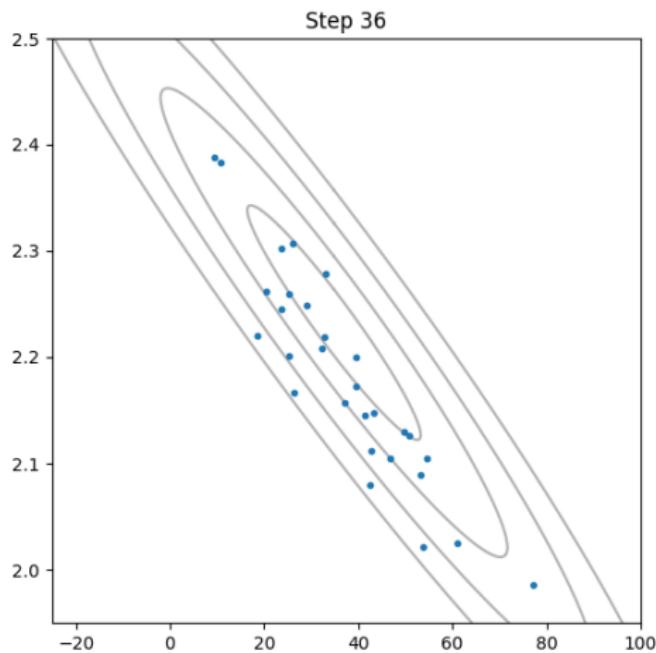
Emcee demo



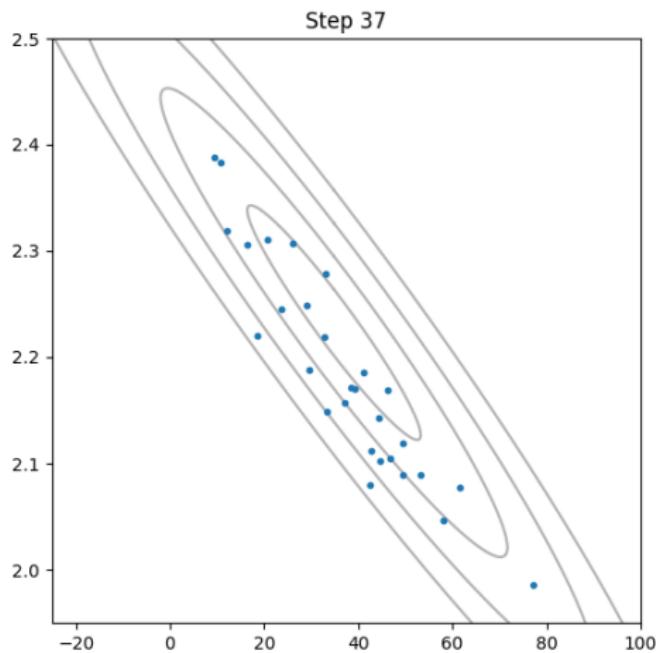
Emcee demo



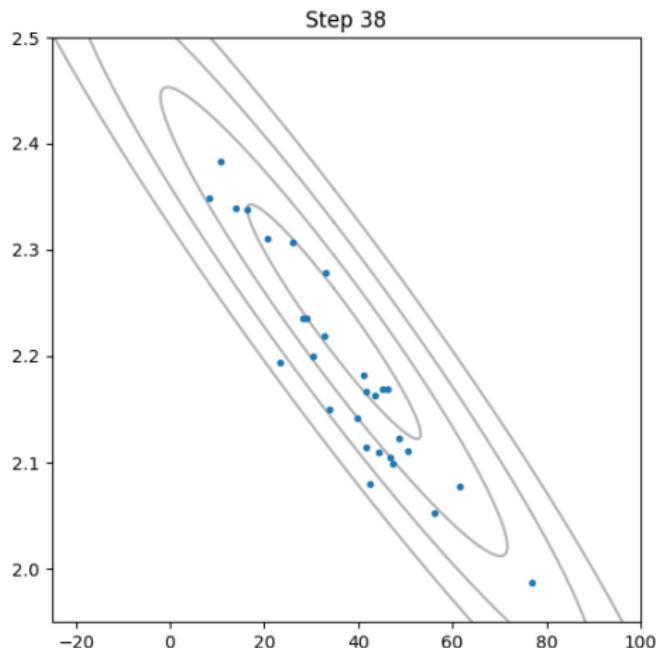
Emcee demo



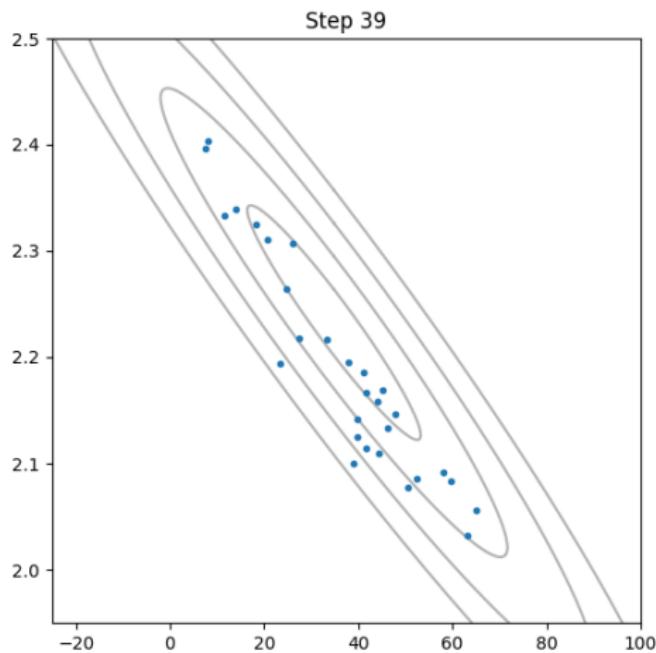
Emcee demo



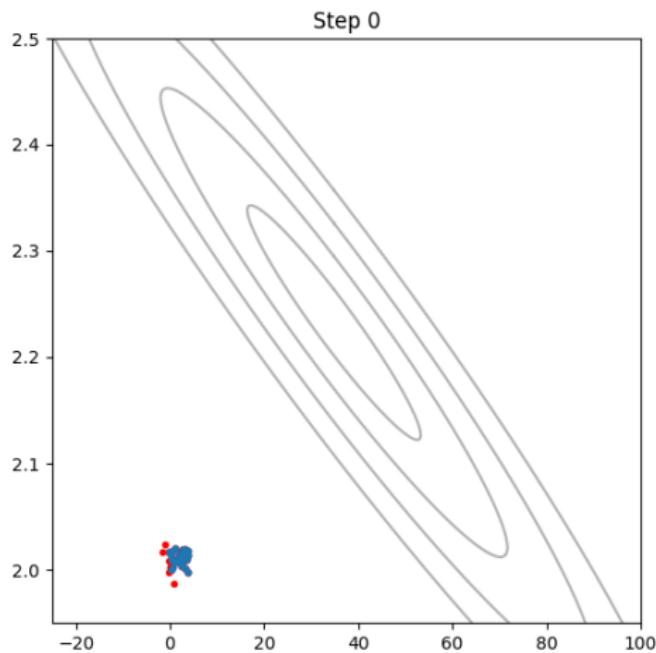
Emcee demo



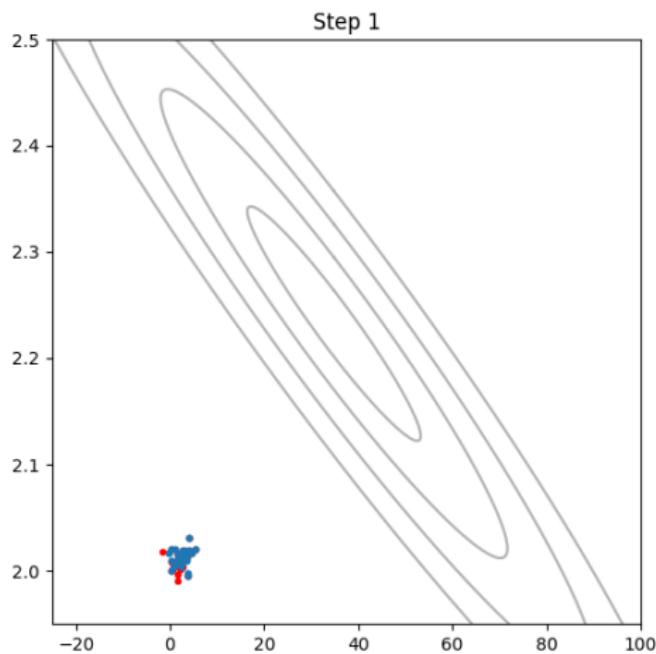
Emcee demo



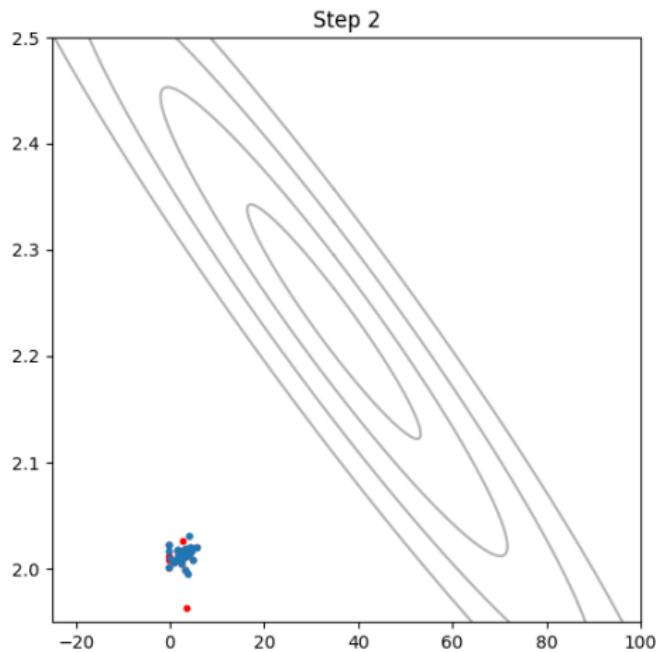
Emcee demo



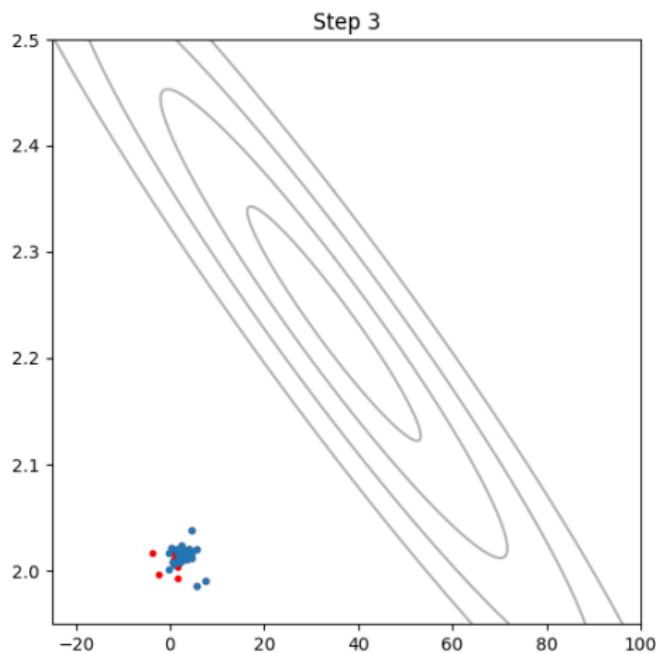
Emcee demo



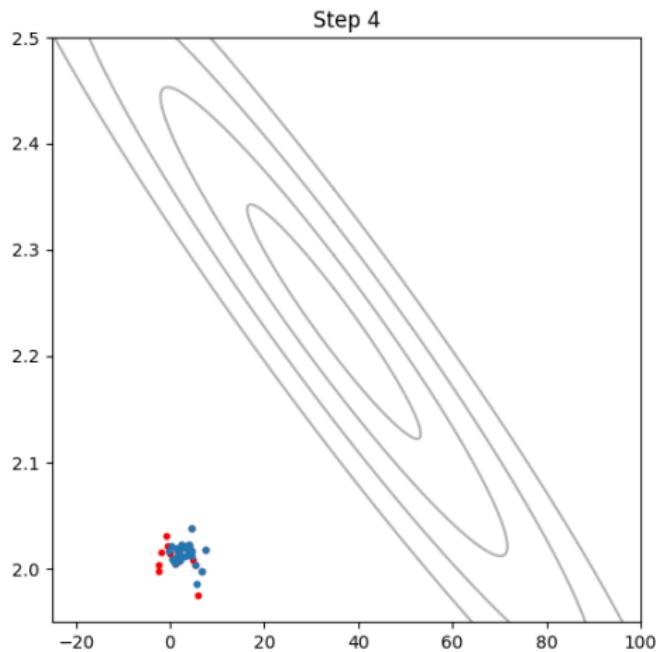
Emcee demo



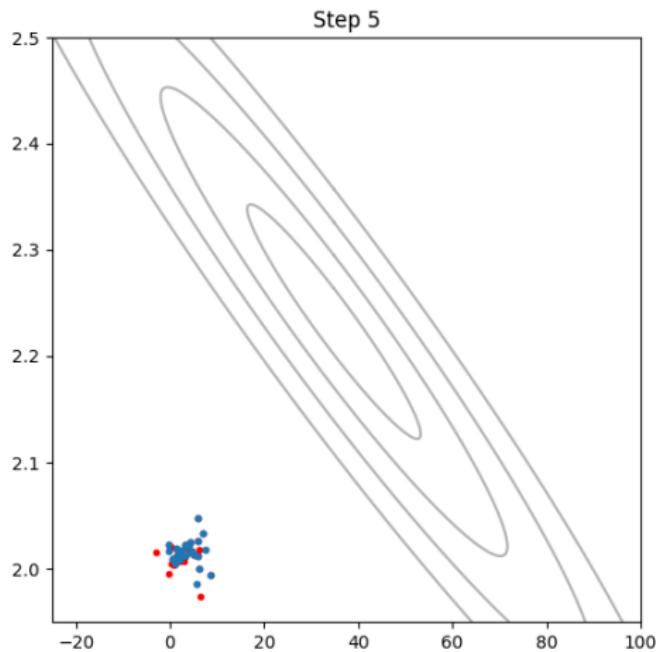
Emcee demo



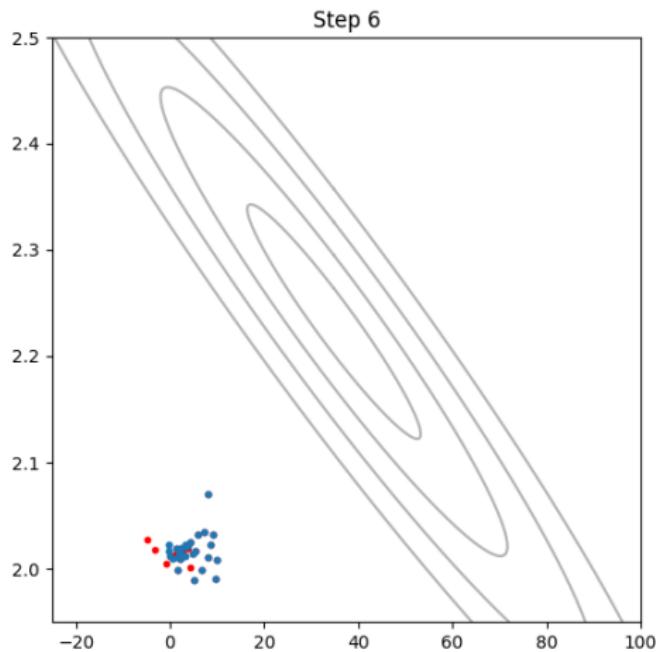
Emcee demo



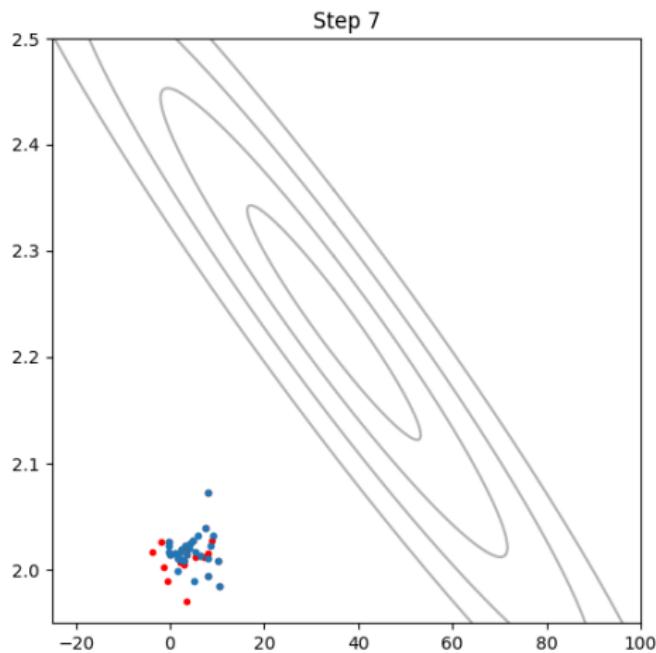
Emcee demo



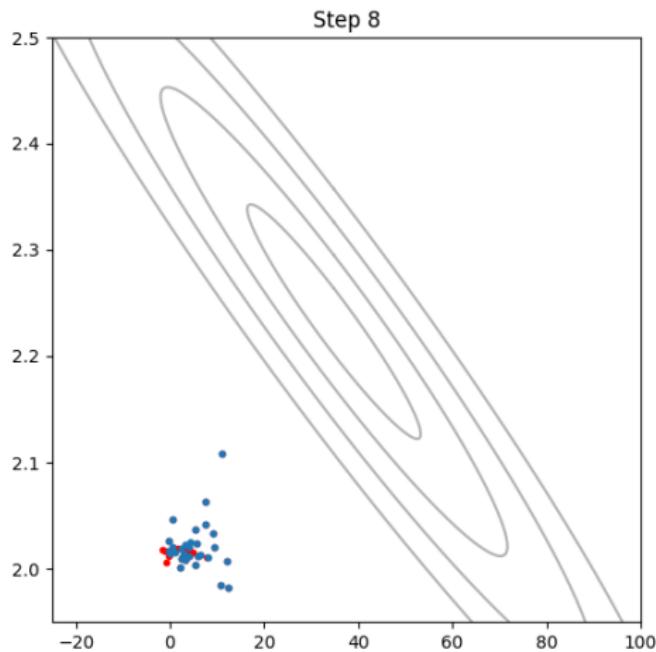
Emcee demo



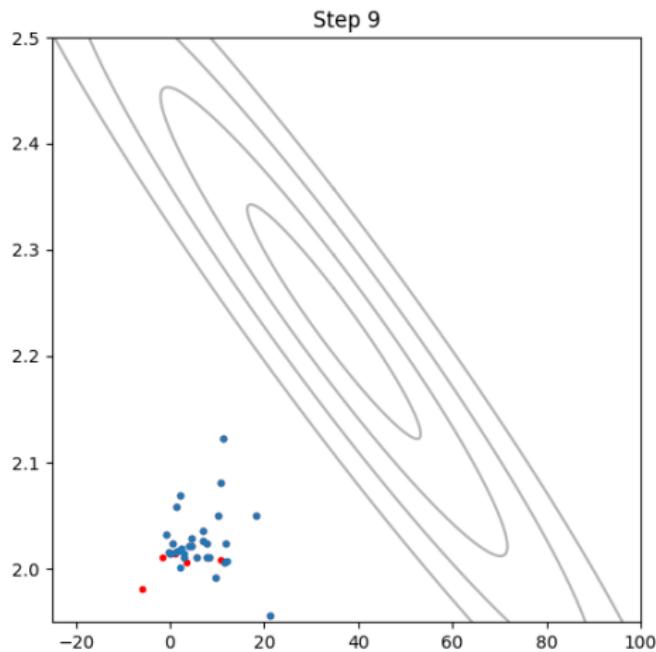
Emcee demo



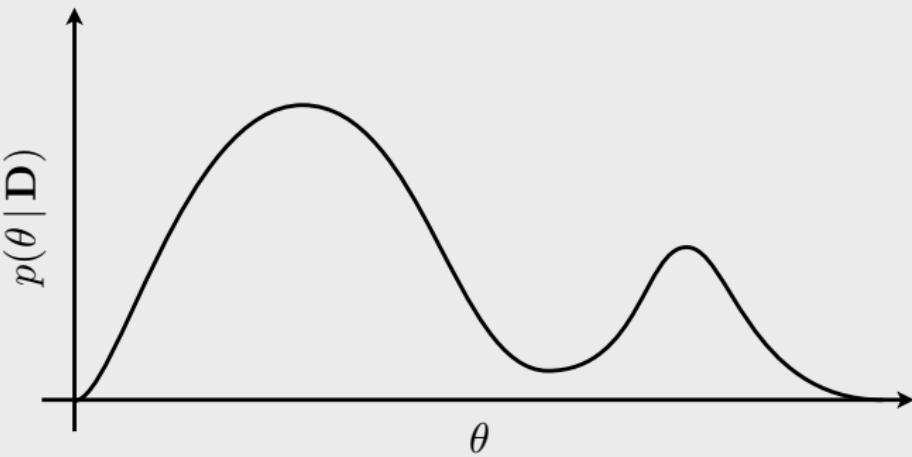
Emcee demo



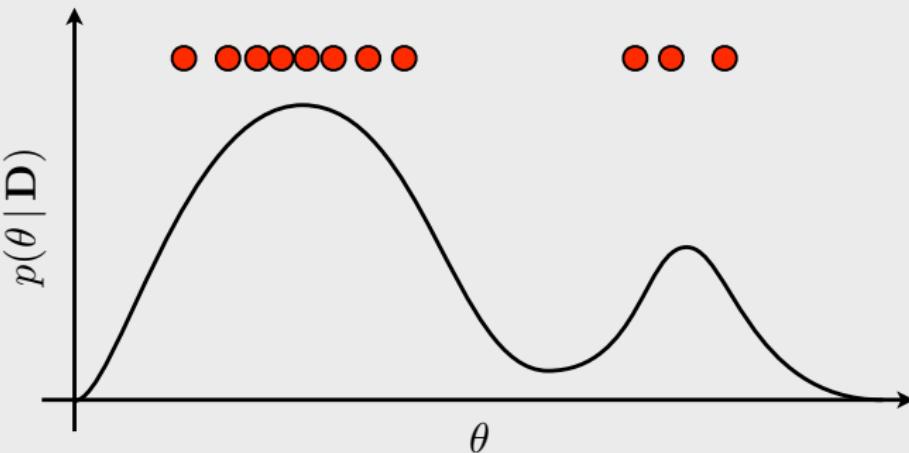
Emcee demo



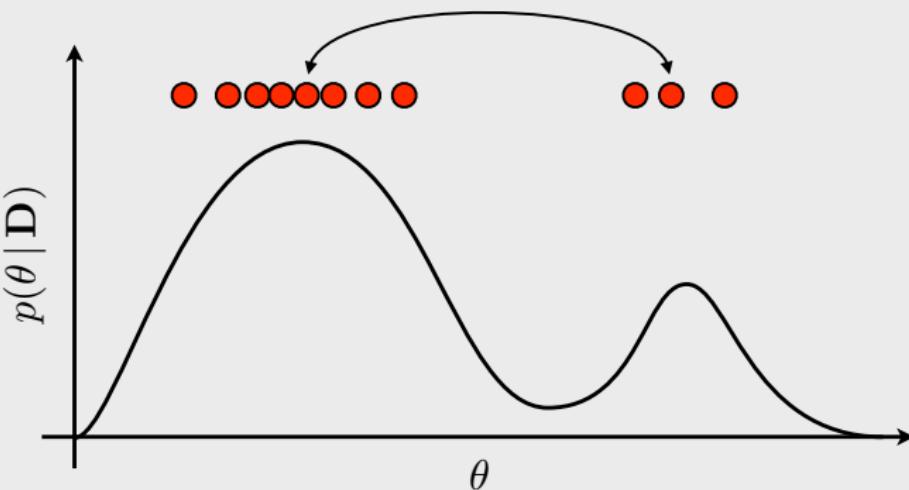
what about multimodal densities?



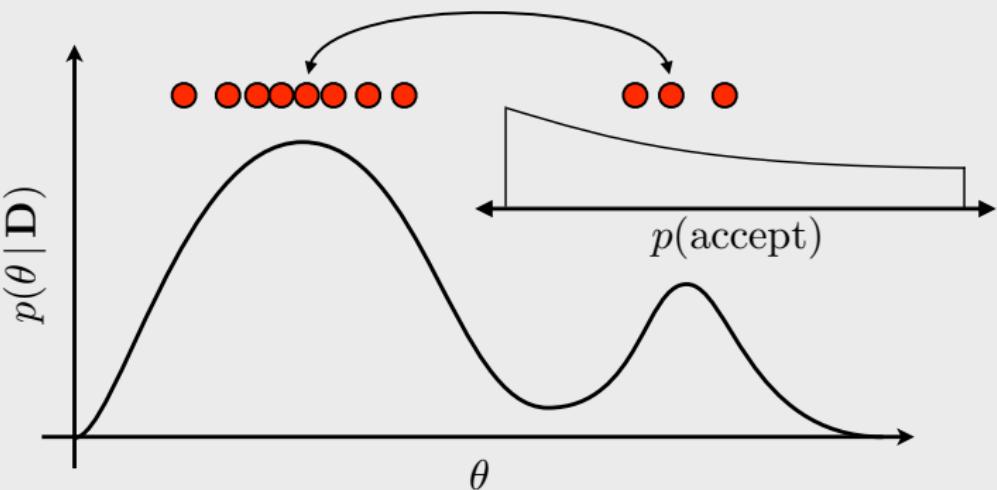
what about multimodal densities?



what about multimodal densities?



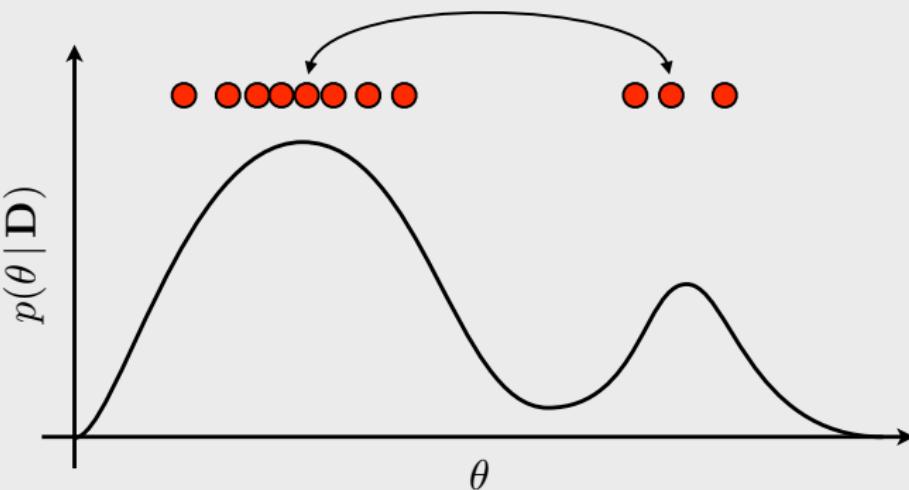
what about multimodal densities?



Differential Evolution move

- ▶ `emcee` allows us to use different *move* types (different *proposal* functions)
- ▶ The **Differential Evolution** (DE) move can improve the sampling for multi-modal distributions
- ▶ DE move: randomly select *two* “helpers”
- ▶ Propose moving by their **vector difference**
- ▶ (If they are from different modes, this proposes *jumping between modes*)
- ▶ Mixing in a fraction of DE moves with the regular “Stretch” move works well!

what about multimodal densities?



Summary

- ▶ Traditional Metropolis–Hastings MCMC suffers from a *lack of affine invariance* – requires *tuning parameters* that change for each specific probability function
- ▶ *Ensemble samplers* like `emcee` use the *distribution of the walkers* to achieve *affine invariance*
- ▶ → much easier to use, and faster sampling
- ▶ (Huge side effect: parallelizable!)
- ▶ Multi-modal distributions still hard, but *DE Move* can help