

# Markov Chain Monte Carlo

Dustin Lang  
Perimeter Institute for Theoretical Physics

PSI Numerical Methods 2024

Borrowing heavily from Dan Foreman-Mackey's slides  
*<https://speakerdeck.com/dfm/data-analysis-with-mcmc>*

These slides are available at  
*<https://github.com/dstndstn/MCMC-talk>*

---


**data analysis** with

# Markov chain Monte Carlo


---

**Dan Foreman-Mackey**

CCPP@NYU



a sketch of  
**The graphical model** of my research.



a sketch of  
**The graphical model** of my research.

$$p(\text{data} \mid \text{physics})$$

likelihood function/generative model

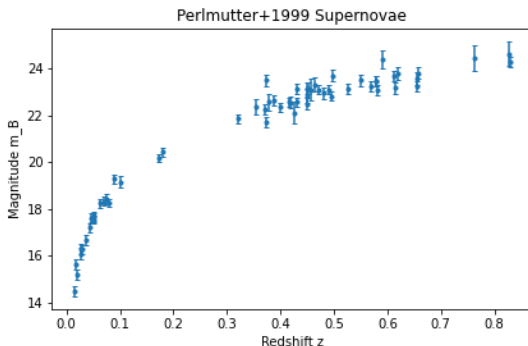


$$p(\text{physics} \mid \text{data}) \propto p(\text{physics}) p(\text{data} \mid \text{physics})$$

posterior probability

## An example

- ▶ Perlmutter+1999 (<https://arxiv.org/abs/astro-ph/9812133>)
- ▶ Measured the observed peak brightnesses of a sample of type-1a supernovae (in astronomer “mag” units), and the redshifts (“z”) of the supernova host galaxies
- ▶  $\text{mag} = \text{mag}_{\text{intrinsic}} + \text{lightness\_distance}(z, \text{parameters})$



# An example

- ▶ Generative model:


$$\text{mag}_i = \text{mag}_{\text{intrinsic}} + \text{luminosity\_distance}(z_i, \text{parameters}) + \epsilon_i$$

- ▶ Probability of a data point given a model (“likelihood”):

$$p(\text{mag}_i | \text{params}) = \text{Gaussian}(\text{mag}_i | \mu = f(z_i, \text{params}), \sigma_i^2)$$


- ▶ Probability of whole data set given a model:

$$p(\{\text{mag}_i\} | \Omega_M, \Omega_\Lambda) = \prod_i \mathcal{N}(\text{mag}_i | \text{mag}_{\text{int}} + D_L(z_i, \Omega_M, \Omega_\Lambda), \sigma_i^2)$$



## An example


- ▶ Use Bayes' theorem to convert data likelihoods into constraints on the model parameters  $\theta = \{\Omega_M, \Omega_\Lambda\}$
- ▶  $p(\theta | \text{data}) \propto p(\theta) p(\text{data} | \theta)$
- ▶  $p(\Omega_M, \Omega_\Lambda | \{\text{mag}_i\}) \propto p(\Omega_M, \Omega_\Lambda) \prod_i \mathcal{N}(\text{mag}_i | \text{mag}_{\text{int}} + D_L(z_i, \Omega_M, \Omega_\Lambda), \sigma_i^2)$





## An example


- ▶ Then they ran MCMC...
- ▶ Resulting parameter constraints (blue ellipse):



# Why we often need MCMC


- ▶ Real-life models and likelihoods are often complex
- ▶ ... so the resulting **constraints** have complicated distributions (not Gaussians!)
- ▶ ... but we can represent them with **samplings**
- ▶ MCMC is used for drawing samples from probability distributions that we can compute numerically but cannot solve analytically

# Samplings to represent constraints - examples



- ▶ From <https://arxiv.org/abs/1910.04899>
- ▶ With a sampling: **Marginalize** over a parameter by projecting it out

# Samplings to represent constraints - examples



► From <https://arxiv.org/abs/1611.00036>

# MCMC


draws samples from a probability function

and all you need to be able to do is


**evaluate**

the function

(up to a constant)




# Metropolis-Hastings





**Metropolis-Hastings**  
in an ideal world

start here  
perhaps




**Metropolis-Hastings**  
in an ideal world






**Metropolis-Hastings**  
in an ideal world



$$p(\text{accept}) = \min \left( 1, \frac{p(\mathbf{x})}{p(\mathbf{x}')} \frac{q(\mathbf{x}; \mathbf{x}')}{q(\mathbf{x}'; \mathbf{x})} \right)$$


**Metropolis-Hastings**  
in an ideal world



$$p(\text{accept}) = \min \left( 1, \frac{p(\mathbf{x})}{p(\mathbf{x}')} \frac{q(\mathbf{x}; \mathbf{x}')}{q(\mathbf{x}'; \mathbf{x})} \right)$$

definitely.

**Metropolis-Hastings**  
in an ideal world



$$\mathbf{x}' \sim q(\mathbf{x}'; \mathbf{x})$$


only **relative** probabilities

$$p(\text{accept}) = \min \left( 1, \frac{p(\mathbf{x})}{p(\mathbf{x}')} \frac{q(\mathbf{x}; \mathbf{x}')}{q(\mathbf{x}'; \mathbf{x})} \right)$$


definitely.

# Metropolis–Hastings

in an ideal world




**Metropolis–Hastings**  
in an ideal world




# Metropolis–Hastings

in an ideal world




**Metropolis-Hastings**  
in an ideal world




not this time.

**Metropolis-Hastings**  
in an ideal world

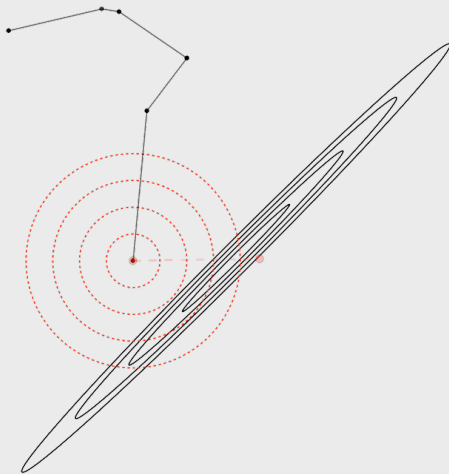





**Metropolis–Hastings**  
in an ideal world



**Metropolis-Hastings**  
in an ideal world



**Metropolis-Hastings**  
in an ideal world



**Metropolis-Hastings**  
in an ideal world

## Note

- ▶ In the previous slides, there is an unfortunate typo! In all cases,

$$p(\text{accept}) = \min\left(1, \frac{p(x)}{p(x')} \frac{q(x; x')}{q(x'; x)}\right)$$

should actually be

$$p(\text{accept}) = \min\left(1, \frac{p(x')}{p(x)} \frac{q(x'; x)}{q(x; x')}\right)$$

That is, **how good is the new place** divided by **how good was the old place?**

## About the name

- ▶ **Monte Carlo**: a reference to the famous Monte Carlo Casino in Monaco, alluding to the randomness used in the algorithm
- ▶ **Markov Chain**: a list of samples, where each one is generated by a process that only looks at the previous one.
- ▶ **Markov**: a 19th-century Russian mathematician and impressive-moustache-haver with an [extensive list of things named after him](#)
- ▶ **Metropolis–Hastings**: lead authors of 1953 and 1970 papers (resp.) giving the algorithm with symmetric and general proposal distributions (resp.)

# The Algorithm (1)

```
function mcmc(prob_func, propose_func, initial_pos, nsteps)
    p = initial_pos
    prob = prob_func(p)
    chain = []
    for i in 1:nsteps
        # propose a new position in parameter space
        # ...
        # compute probability at new position
        # ...
        # decide whether to jump to the new position
        if # ...
            # ...
            # ...
        end
        # save the position
        append!(chain, p)
    end
    return chain
end
```

## The Algorithm (2)

```
function mcmc(prob_func, propose_func, initial_pos, nsteps)
    p = initial_pos
    prob = prob_func(p)
    chain = []
    for i in 1:nsteps
        # propose a new position in parameter space
        p_new = propose_func(p)
        # compute probability at new position
        prob_new = prob_func(p_new)
        # decide whether to jump to the new position
        if prob_new / prob > uniform_random()
            p = p_new
            prob = prob_new
        end
        # save the position
        append!(chain, p)
    end
    return chain
end
```



## The Algorithm (3)


```
function mcmc(logprob_func, propose_func, initial_pos, nsteps)
    p = initial_pos
    logprob = logprob_func(p)
    chain = []
    for i in 1:nsteps
        # propose a new position in parameter space
        p_new = propose_func(p)
        # compute probability at new position
        logprob_new = logprob_func(p_new)
        # decide whether to jump to the new position
        if exp(logprob_new - logprob) > uniform_random()
            p = p_new
            logprob = logprob_new
        end
        # save the position
        append!(chain, p)
    end
    return chain
end
```

## The Algorithm (4)

```
function mcmc(logprob_func, propose_func, initial_pos, nsteps)
    p = initial_pos
    logprob = logprob_func(p)
    chain = []
    naccept = 0
    for i in 1:nsteps
        # propose a new position in parameter space
        p_new = propose_func(p)
        # compute probability at new position
        logprob_new = logprob_func(p_new)
        # decide whether to jump to the new position
        if exp(logprob_new - logprob) > uniform_random():
            p = p_new
            logprob = logprob_new
            naccept += 1
        end
        # save the position
        append!(chain, p)
    end
    return chain, naccept/nsteps
```


# Practicalities

- ▶ How do I choose a proposal distribution?
- ▶ How many steps do I have to take?




# Metropolis-Hastings

in the real world




**Metropolis-Hastings**  
in the real world




# Metropolis-Hastings

in the real world




**Metropolis-Hastings**  
in the real world




**Metropolis-Hastings**  
in the real world






the  
**Small Acceptance Fraction**  
problem

**Metropolis–Hastings**  
in the real world




**Metropolis–Hastings**  
in the real world



**Metropolis–Hastings**  
in the real world




**Metropolis–Hastings**  
in the real world




**Metropolis-Hastings**  
in the real world

that's  $O(D+1)/2$   
tuning  
parameters




**Metropolis–Hastings**  
in the real world

that's  $D(D+1)/2$  tuning parameters  
the dimension of your problem



**Metropolis–Hastings**  
in the real world




# Metropolis-Hastings

in the real world



## A connection to symmetries

- ▶ In Metropolis–Hastings MCMC, the *proposal distribution* needs **tuning parameters**, especially as dimensionality increases
- ▶ Can be seen as a lack of **symmetry** in the algorithm—the algorithm is sensitive to the parameterization of the problem
- ▶ For example, it's not invariant to an **affine** transformation
- ▶ **Next lecture**, I'll show you an alternative algorithm that **does** have affine invariance



# How many samples do I need?

- ▶ Burn-in — skip the first  $N$  samples
- ▶ *Has my chain converged?*
- ▶ MCMC produces **correlated** samples, so
  - ▶ How correlated are my samples?
    - ▶ Can measure the *autocorrelation time*  $\tau$
    - ▶ Keep  $1/\tau$  of the MCMC samples
    - ▶ eg <https://github.com/dfm/acor>
  - ▶ How many uncorrelated samples do I need?
    - ▶ No easy general answer to this question!
    - ▶ “How many can you afford?”

# How many samples do I need?

- ▶ Burn-in — skip the first  $N$  samples
- ▶ *Has my chain converged?*
- ▶ MCMC produces **correlated** samples, so
  - ▶ How correlated are my samples?
    - ▶ Can measure the *autocorrelation time*  $\tau$
    - ▶ Keep  $1/\tau$  of the MCMC samples
    - ▶ eg <https://github.com/dfm/acor>
  - ▶ How many uncorrelated samples do I need?
    - ▶ No easy general answer to this question!
    - ▶ “How many can you afford?”

# How many samples do I need?

- ▶ Burn-in — skip the first  $N$  samples
- ▶ *Has my chain converged?*
- ▶ MCMC produces **correlated** samples, so
  - ▶ How correlated are my samples?
    - ▶ Can measure the *autocorrelation time*  $\tau$
    - ▶ Keep  $1/\tau$  of the MCMC samples
    - ▶ eg <https://github.com/dfm/acor>
  - ▶ How many uncorrelated samples do I need?
    - ▶ No easy general answer to this question!
    - ▶ “How many can you afford?”

# Conclusions

- ▶ MCMC remains an essential tool for probabilistic inference
- ▶ For science: lets us constrain model parameters based on data (Bayesian inference)
- ▶ Beguilingly simple algorithm, but difficult practicalities
- ▶ MCMC has beautiful theoretical guarantees... as compute time  $\rightarrow \infty$